Resilient Grid Intelligence: A Standards-Guided Graph Neural Network Framework for Incipient Fault Prediction and Cyber-Physical Anomaly Discrimination in Saudi Medium-Voltage Distribution Networks

Moustafa Elbehairy[1,2], Ahmad Abuseeni[1], Nasser Al-Maawi[3], Amir Raouf[4], and Ahmed Abdulrahman Al-Jafin[5]
[1]Nukhbaat Hulul Alamal, Riyadh, Saudi Arabia [2]SystechHub, Germany [3]United Specialists Co., Saudi Arabia [4]Karaz Platform, Saudi Arabia [5]Energy Towers Company for Engineering, Saudi Arabia
m.elbehairy@nukhbaa.sa, a.abuseen@nukhbaa.sa, nasser.a@usco-sa.com, amir.raouf@outlook.sa, aaljafin@etco.com.sa | Corresponding author: Moustafa Elbehairy

Abstract—Most customer interruptions originate in medium-voltage networks, and distribution transformers and cable joints emit precursors days before failure. VoltyGuard represents a 13.8 kV-class feeder as a dynamic graph and fuses standards-based diagnostics, physics-based thermal and power-flow models and a spatio-temporal graph neural network, with alerts gated by agreement between the standards and learned layers. A cyber-physical discriminator separates incipient faults from false data injection and feeds an explicit decision policy; a regulation-aware scheduler maps lead time to actions admissible under the Saudi Electricity Regulatory Authority's standards. On 24 simulated Riyadh summers of the IEEE 33-bus feeder, with thresholds and the false-alarm budget fixed on validation data, the framework detects 82% of failures at the standards alarm's budget (alarm: 85%) with a median warning of 7.2 days, separates faults from attacks in 90% of episodes, and schedules 21 of 27 failures as planned outages while holding suspected attacks for security review.

Index Terms—Cyber-physical security, dissolved gas analysis, distribution reliability, explainable AI, false data injection, graph neural networks, partial discharge, predictive maintenance, Saudi Arabia, transformer thermal model.

I. Introduction

Saudi distribution networks are modernising quickly. Distribution automation was reported at about 40% by the end of 2025, with 12.3 GW of renewable capacity connected to the national grid [1], and the Saudi Electricity Company reported improvements of 17% in SAIDI and 19% in SAIFI for 2024 [2]. Sustaining those gains means moving from restoration after an outage to anticipation before it, and the medium-voltage (MV) network is where that shift pays most, because it is where most customer-minutes lost originate. The regulatory frame makes the economics explicit. The guaranteed service standards of the Saudi Electricity Regulatory Authority (SERA) require at least two days of notice before a planned interruption, restoration within six hours of a planned outage and within three hours of an unplanned one, with automatic customer compensation when a limit is breached [3]. A warning that arrives earlier than the 48-hour notice period can therefore be converted into a planned, compensated-free outage; one that arrives later cannot.

Two obstacles have slowed the use of machine learning for this purpose in utility control rooms. The first is trust: engineers do not act on opaque scores that contradict diagnostic standards they have applied for decades, a concern that the interpretability literature shares for high-stakes decisions [4]. The second is security. A model that consumes operational data inherits the attack surface of that data: coordinated false data injection (FDI) can pass residual-based bad-data detection [5] and either provoke an unnecessary switching operation or hide a genuine degradation. An early-warning system that cannot tell a failing transformer from a spoofed sensor is, in the attacker's hands, a remote switching tool.

This paper reports the full study announced in our accepted extended abstract. We make four contributions. (1) A standards-guided hybrid architecture in which a graph neural network (GNN) is fed with, and gated by, the outputs of IEEE C57.104 [6], IEC 60599 [7], IEC 60270 [8] and the IEC 60076-7 thermal model [9]. (2) A cyber-physical discriminator that uses cross-modal and power-flow consistency across neighbouring nodes to separate incipient faults, data-integrity attacks and sensor faults. (3) A regulation-aware scheduler that maps warning lead time to actions admissible under SERA's guaranteed standards, including an N-1 load-transfer check in a power-flow model. (4) A reproducible evaluation on the IEEE 33-bus feeder [10] under Riyadh summer load and temperature, with 116 degradation episodes, 75 injection attacks and 48 sensor malfunctions across 24 independent realisations, against static-threshold and per-asset learning baselines. Code, data generator and results are public.

II. Background and Related Work

Transformer condition assessment is mature and standardised. Dissolved gas analysis (DGA) is interpreted with the gas-by-gas screening values and rate-of-change limits of IEEE C57.104-2019 [6] and the ratio methods of IEC 60599, including the Duval triangle [7], [11]; partial discharge (PD) is quantified per IEC 60270 [8]; and thermal ageing follows the loading guide IEC 60076-7 [9], whose exponential top-oil and hot-spot model has been validated for distribution units [12]. For underground cables, failures of joints and terminations are preceded by sub-cycle self-clearing arcing events that recur with increasing frequency; their detection and location from feeder-head recordings is well documented [13]–[15]. These methods, however, are applied asset by asset. A transformer is assessed without reference to the joint feeding it, and a recurring transient is attributed to a segment without reference to the voltage profile its neighbours observe.

Graph neural networks exploit grid topology directly through message passing [16], [17] and have been reviewed

for power-system use [18]. They have located faults in distribution systems [19], localised faulted lines from sparse PMU data [20], estimated unmeasured states [21] and detected stealthy FDI attacks [22]. Machine learning has also been used to predict distribution faults under heatwaves from weather and asset data [23]. To our knowledge these strands remain separate: fault-prediction models assume trustworthy data, and FDI detectors do not reason about physical degradation. Neither connects its output to what a regulated operator may actually do with it. The gap this paper addresses is a single, standards-anchored model that separates normal operation, incipient physical faults and data-integrity attacks on the same graph, and that converts its output into actions admissible under national service regulation.

III. Proposed Framework

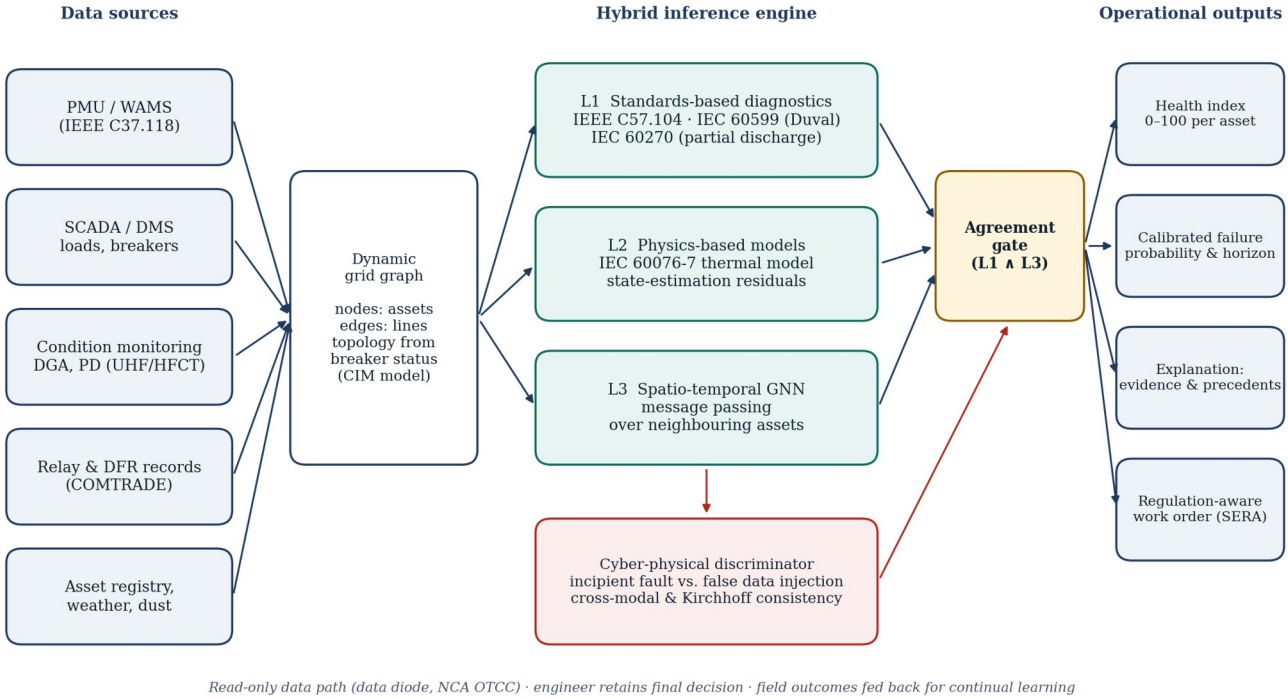


Fig. 1. VoltyGuard framework: multi-source data on a dynamic grid graph, three-layer hybrid inference, cyber-physical discrimination and agreement-gated operational outputs.

A. Dynamic Grid Graph

The feeder is a graph whose nodes are assets (primary-substation busbar, cable sections and joints, distribution transformers, measurement points) and whose edges are electrical connections carrying series impedance. Topology, including normally-open points, is imported from the utility network model and updated from switch status, so reconfiguration changes the adjacency without retraining. Node features combine SCADA loads and voltages, branch flows, protection and fault-recorder events, condition-monitoring data where available (dissolved gases, PD, top-oil temperature), upstream synchrophasor data where available [24], asset rating and age, and ambient temperature. Fig. 1 summarises the architecture.

B. Three-Layer Hybrid Inference

Layer 1 applies the standards-based rules of Section II to each asset and emits an interpretable code (normal, caution, alarm): DGA status from the C57.104 screening and rate limits, PD caution and alarm levels, self-clearing fault counts, and top-oil and hot-spot limits. *Layer 2* computes physics-based quantities from the measured data: the IEC 60076-7 top-oil and hot-spot temperatures and the relative ageing rate implied by the measured load and ambient temperature, the residual between measured and modelled top-oil temperature, and two power-flow residuals obtained by solving the network with the measured loads, namely the voltage residual at each bus and the branch power-balance residual (measured flow minus the flow implied by all measured loads). *Layer 3* is a spatio-temporal GNN. A gated recurrent unit encodes the last 24 h of every node's features; two GraphSAGE-style layers [17] then exchange messages over the admittance-weighted adjacency, so that each asset's representation includes the state of its electrical neighbours; two heads output the probability of failure within seven days and the probability of each anomaly class. Probabilities are calibrated by temperature scaling on validation data [25]. A *critical* alert requires agreement between Layer 1 (code at least caution within the previous 24 h) and Layer 3 (calibrated probability above a threshold for three consecutive hours); a Layer-3 alert without Layer-1 support is routed to engineer review with the evidence of both layers.

C. Cyber-Physical Anomaly Discrimination

Genuine incipient faults leave physically consistent signatures across independent modalities and neighbouring nodes: rising gases corroborate hot-spot temperatures computed from load, PD activity corroborates self-clearing events, and electrical quantities obey the power-flow equations across the graph. An attacker who manipulates one channel

must spoof several modalities and several nodes consistently to avoid detection. The discriminator is the second head of Layer 3, trained on the Layer-2 residuals and on the neighbours' features, and it classifies each anomaly as incipient fault, suspected data-integrity attack or sensor malfunction. The class output is advisory and feeds an explicit decision policy: each alert cluster is routed by its majority class, incipient fault or unclassified to the maintenance scheduler, suspected attack to security review without any switching recommendation, and sensor malfunction to an instrumentation check. Section V-C evaluates this policy, including the case of genuine degradation under a masking attack, which stays visible because the attacker does not control the hot-spot temperature implied by load or the neighbours' voltages.

D. Regulation-Aware Maintenance Scheduler

The scheduler converts each critical alert into an admissible action given its lead time, the N-1 margin of adjacent assets and customer criticality (Table I). For a cable-section alert the N-1 check opens the section, closes the best normally-open tie and solves the power flow, accepting the transfer when all bus voltages stay above 0.95 p.u. and no cable exceeds its ampacity. For a transformer alert the check compares the spare capacity of the adjacent unit with the load to be transferred.

TABLE I
MAPPING OF PREDICTION LEAD TIME TO ADMISSIBLE ACTIONS

Lead time	**Admissible action proposed by the scheduler**	**SERA guaranteed standard**
≥ 48 h	Planned outage with prior customer notice; repair or replacement in a low-impact window	Std. 4 (notice ≥ 2 days); Std. 5 (restore ≤ 6 h)
< 48 h, N-1 margin	Load transfer to adjacent feeder or transformer; no customer interruption	—
< 48 h, no margin	Pre-positioned crew, spare and mobile substation; controlled emergency response	Std. 6 (restore ≤ 3 h)
Fault occurred	Graph-based fault localisation to shorten restoration	Std. 6 (restore ≤ 3 h)

IV. Simulation Study Design

A. Feeder, Climate and Load

We use the IEEE 33-bus radial feeder of Baran and Wu [10] as shipped with pandapower [26] (12.66 kV, 3.715 MW and 2.3 MVAr at peak, 32 cable sections and five normally-open ties), representative of the 13.8 kV class operated in Saudi cities. Each load bus carries a distribution transformer sized at 0.95–1.35 times its peak apparent load; cable ampacities are assigned from the peak-case currents. Each realisation covers one summer of 92 days at hourly resolution. Ambient temperature follows Riyadh June–August normals (daily mean about 36 °C, a 16 K diurnal range with the maximum at 15:00, day-to-day autoregressive variation and two to three heatwave episodes of three to six days adding 3–5 K; on average 10.6 heatwave days per summer). Load combines residential and commercial diurnal shapes with a cooling sensitivity of 2.5% per kelvin above 30 °C, a weekend reduction and lognormal noise; the hourly network state is obtained by a backward-forward-sweep power flow.

B. Degradation, Attack and Sensor-Fault Scenarios

Three transformer and joint processes are simulated per realisation (two to three of each). *Transformer thermal ageing*: under-sized units (rated at 0.68–0.84 times peak load) are driven through the IEC 60076-7 model (distribution ONAN parameters: top-oil rise 55 K, hot-spot gradient 23 K, loss ratio 5, exponents 0.8 and 1.6, oil time constant 180 min); the relative ageing rate accumulates into a damage index normalised so that failure occurs at a random day of the summer, and thermal-fault gases (C_2H_4, CH_4, C_2H_6, H_2, with C_2H_2 only near the end) are generated at a rate that rises with hot-spot temperature and accelerates once the damage index exceeds one half. *Cable-joint degradation*: a latent damage index grows from a random onset over 15–38 days, accelerated by load current; it drives a Poisson process of self-clearing faults (up to about four per day near failure, attributed to the correct segment with 80% probability and otherwise to an adjacent one, to mimic location error), a rising PD level where a monitor is installed, and a growing contact resistance that is written into the power-flow model so that neighbouring voltages respond physically. *FDI attacks* of three kinds target healthy or degrading assets for two to eight days: a *provoke* attack inflates the top-oil sensor and the online DGA of a healthy transformer to imitate degradation; a *stealthy electrical* attack scales the load and the feeding-branch measurements of a node consistently with the local voltage drop, so that the node's own power-flow residual stays small while the parent branch's balance is violated; and a *mask* attack replays the condition data of a degrading asset from 15 days earlier during its final days. *Sensor malfunctions* are a stuck or drifting top-oil sensor. Measurement noise is 0.3% on voltages, 1% on powers and currents, 1 K on temperatures and 5% on gas concentrations. Online DGA is available on 51% of transformers (daily samples; the others are sampled monthly) and PD monitors on 48% of joints, which exposes the models to partial observability. Table II summarises the data.

TABLE II
SIMULATED DATA SET (24 REALISATIONS × 92 DAYS × 33 NODES)

Quantity	**Value**
Node-hours (train / validation / test realisations)	1,748,736 (14 / 4 / 6 realisations)
Features per node (raw + rolling 24 h / 168 h statistics)	30 + 52
Node-hours labelled incipient / attack / sensor fault	57,958 / 8,508 / 4,614
Failure episodes: cable joints / transformers (test)	59 / 57 (14 / 13)
FDI episodes: provoke / electrical / mask (test)	38 / 24 / 13 (11 / 6 / 4)
Sensor-fault episodes: stuck / drift (test)	24 / 24 (6 / 6)
Layer-1 caution-or-alarm rate on normal node-hours	3.6%
Simulation time per realisation (laptop CPU)	25 s

C. Models, Baselines and Ablations

Static rules use the Layer-1 code alone (alarm, or caution-or-alarm). The *per-asset baseline* is a gradient-boosted tree ensemble [27] trained on each node's own features and rolling statistics, without neighbours; it represents the asset-isolated learning approach of [23]. *VoltyGuard* is the network of Section III-B with hidden size 64 and 21,173 parameters, trained with class-weighted losses, AdamW, early stopping on the validation realisations and 5 random seeds (three for each ablation, which were trained first; like-for-like comparisons

use the first three seeds of the full model). Two ablations remove the graph (adjacency set to zero, so each node sees only itself) and remove the Layer-1/2 features (the network sees raw measurements only). Realisations are split 14/4/6 into training, validation and test, so that no test asset, episode or weather trace is seen in training.

D. Metrics

An alert is raised when the calibrated failure probability exceeds a threshold for three consecutive hours; the threshold is chosen on the validation realisations to maximise episode-level F1 (the *F1 operating point*); a second, *matched-budget* threshold is the lowest one whose validation false-cluster rate stays within the static alarm's false-cluster rate on the same validation realisations (0.25 per 100 asset-days). Both thresholds and the budget are therefore fixed on validation data before any test metric is read, and the test false-cluster rate actually realised at the matched threshold is reported rather than assumed equal. A failure episode counts as detected if an alert occurs on that asset within the 45 days before failure, and its *warning lead time* is the interval from the first such alert to failure. Alert runs on assets that do not fail within 45 days are *false clusters*, reported per 100 asset-days. Episode recall and precision follow. *Discrimination accuracy* is the share of anomaly episodes whose majority class over the episode is correct; the fault-versus-attack figure restricts this to incipient-fault and FDI episodes. Calibration is the expected calibration error (ECE) on test node-hours [25]. The SERA mapping of Table I is applied to every detected failure at the hour of its first alert, and again under the decision policy of Section III-C, at the hour of the first alert routed to the scheduler. For the policy we count attack episodes whose alerts reached the scheduler (unnecessary switching recommendations), failures whose alerts were held as suspected attacks (missed or delayed by the policy) and masked assets that failed while under attack. Recall is reported per seed and per fault class, with a 95% bootstrap interval over test episodes.

V. Results

A. Detection and Warning Lead Time

TABLE III
EPISODE-LEVEL PERFORMANCE ON THE SIX TEST REALISATIONS (27 FAILURES; MEANS OVER SEEDS: 5 FOR VOLTYGUARD, THREE OTHERWISE). COLUMNS 2–6: F1 OPERATING POINT; 7–9: MATCHED FALSE-ALARM BUDGET FIXED ON VALIDATION DATA (0.25 FALSE CLUSTERS / 100 ASSET-DAYS)

Model	Recall	Prec.	FA/ 100	Lead (d)	≥48 h	Rec.@ B	FA@ B	Lead @B
Static alarm (Layer 1 = alarm)	0.85	0.19	0.54	4.0	0.78	—	—	—
Static caution-or-alarm	0.96	0.10	1.31	31.2	1.00	—	—	—
Per-asset GBT	0.74	0.57	0.08	5.9	0.85	0.74	0.05	6.3
GNN without standards/physics	0.74	0.75	0.04	6.5	0.98	0.83	0.18	10.4
GNN without graph	0.63	0.79	0.03	6.2	1.00	0.80	0.10	8.1
VoltyGuard (Layer 1 + 2 + 3)	0.70	0.77	0.04	6.7	0.98	0.82	0.10	7.2
VoltyGuard, agreement-gated	0.70	0.79	0.04	5.6	0.98	0.90	0.20	6.4

FA/100: false alert clusters per 100 asset-days. Rec.@B, FA@B, Lead@B: recall, realised test false-cluster rate and median lead (days) at the matched budget.

TABLE IV
VOLTYGUARD (GATED) BY TRAINING SEED: F1-POINT THRESHOLD AND RECALL BY FAULT CLASS (27 TEST FAILURES: 14 JOINTS, 13 TRANSFORMERS)

Seed	τ	Joints	Transf.	Recall	FA/100	Rec.@B	FA@B
0	0.85	0.93	0.69	0.81	0.03	0.93	0.08
1	0.90	0.93	0.69	0.81	0.01	0.89	0.06
2	0.95	0.93	0.00	0.48	0.01	0.85	0.09
3	0.90	0.93	0.23	0.59	0.03	0.93	0.49
4	0.90	0.93	0.69	0.81	0.11	0.93	0.26

95% bootstrap interval of mean recall over test episodes: 0.56–0.83.

Table III gives the main result. The standards alarm alone detects 85% of the 27 test failures with a median lead of 4.0 days, but only 19% of its alert clusters precede a failure (0.54 false clusters per 100 asset-days, about 63 per feeder-year); the caution level raises recall to 96% at 1.31 false clusters per 100 asset-days. At its F1 operating point VoltyGuard is the most selective model: 77% precision and 0.04 false clusters per 100 asset-days, with 70% recall (seed-to-seed s.d. 14%; bootstrap interval 56%–83%) and a median lead of 6.7 days. The fairer comparison is at equal nuisance. When every learned model is given the standards alarm's own false-alarm budget, VoltyGuard detects 82% of failures with a median lead of 7.2 days, against 74% and 6.3 days for the per-asset ensemble and 85% and 4.0 days for the alarm itself; every VoltyGuard warning at either operating point arrives before the 48-hour notice period. Agreement gating removes the few ungrounded alerts (false clusters 0.04 per 100 asset-days) and shortens the median lead to 5.6 days, because a critical alert now waits for the first standards caution. Lead times differ by fault class (Fig. 3): transformers, whose thermal history is observable through the hot-spot model and whose gases accumulate over weeks, are warned a median 7.5 days ahead at the F1 point, while cable joints, half of which have no PD monitor, give 7.9 days. The static alarm warns transformers earlier still (median 25.0 days), because the hot-spot limit of IEC 60076-7 is exceeded weeks before gassing becomes decisive, but it gives joints only 2.6 days; the learned model's seven-day target horizon concentrates its warnings in the week before failure, and a longer horizon would move them earlier at a cost in precision.

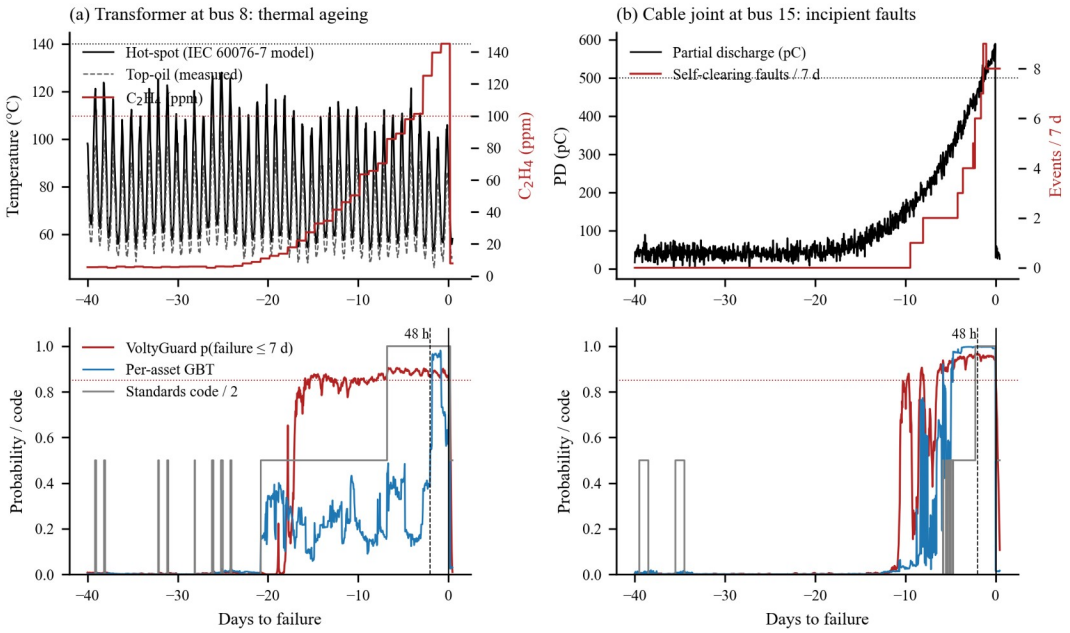


Fig. 2. Two test episodes. Top: precursors of (a) a transformer thermal failure and (b) a cable-joint failure. Bottom: calibrated failure probability of VoltyGuard and of the per-asset model, and the Layer-1 standards code; the dashed line marks the SERA 48-hour notice period.

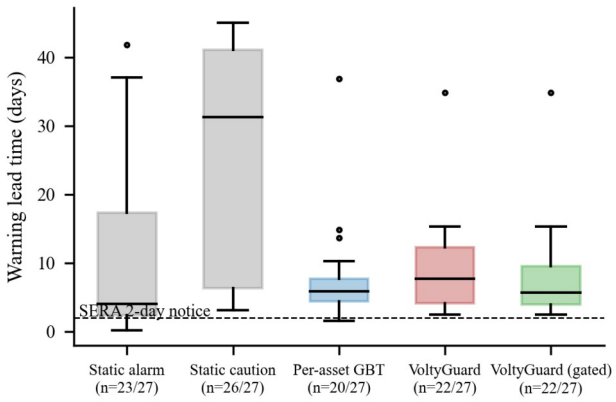


Fig. 3. Warning lead time of detected failures on the test realisations (seed 0). Boxes show quartiles; the dashed line is the SERA two-day notice requirement.

The seed-level figures behind these means are not uniform (Table IV). Joint failures are detected at 0.93 or better by every seed, but transformer recall at the F1 point ranges from 0.00 to 0.69, and seeds 2 and 3 detect fewer than half of the 13 transformer failures. Seed 2 is the extreme case: its validation-optimal threshold is 0.95, the top of the search grid, and at that threshold it detects none of them, whereas one grid step lower (0.90) the same network detects 0.31 of them at 0.06 false clusters per 100 asset-days. The weakness is one of threshold selection on four validation realisations as much as of the representation: at the matched budget, which fixes the threshold by the false-alarm rate instead of by F1, seed 2 reaches 0.85 and the five-seed mean is 0.90. We report the weak seeds rather than exclude them, and we read the F1-point recall as uncertain to the width of its bootstrap interval.

B. Fault-versus-Attack Discrimination

Fig. 4(a) shows the episode-level confusion of VoltyGuard's class head over the 59 test anomaly episodes. Incipient faults and FDI attacks are separated in 91% of the episodes (mean 90% over seeds); 0 of the 21 attacks were read as faults, and three-class accuracy including sensor faults is 86%. The mask attack, which replays a degrading asset's own earlier condition data, is the hardest case: 3 of 4 were flagged as attacks, the others being exposed only by the physical channels the attacker did not control (hot-spot temperature implied by load, neighbouring voltages). The per-asset ensemble, which cannot see the parent branch's power balance, separates faults from attacks in 87% of episodes.

C. Cyber-Aware Decision Policy

Routing every alert cluster by its class turns the discriminator into an operational safeguard that can be scored. Over the 5 seeds of the gated model, 0.2 of the 21 test attack episodes produced failure alerts on the attacked asset; of these, 0.2 were held for security review and 0.0 reached the scheduler as unnecessary switching recommendations (seed 0: 0 of 1). The cost on the physical side is small: 0.2 genuine failures per seed had every alert held as a suspected attack and 0.2 were delayed to a later, fault-classed alert, so policy recall is 70% against 70% without the policy, with a median lead of 5.6 days. The simultaneous fault-plus-attack case is the weak point: of the 4 test assets that were degrading towards failure while their condition data were replayed, the failure was detected and scheduled in time in only 0.6 cases per seed (seed 0: 1 of 4). The hot-spot temperature implied by load and the neighbouring voltages, which the attacker does not control, expose the mask at episode level (Section V-B), but they did not by themselves lift the failure probability above the alert threshold before failure. are not under the attacker's control. False clusters outside any episode split 6.4 to the scheduler and 0.0 to security review per seed. The claim the paper makes is accordingly bounded: the class head is an advisory filter whose residual error is measured here, not a guarantee that no attack can reach a switching decision.

D. Calibration

After temperature scaling the ECE of VoltyGuard's seven-day failure probability is 0.023 ± 0.006 on 437,184 test node-hours (Fig. 4(b)), against 0.007 for the per-asset model. Calibration matters for the scheduler, which reads the probability as a decision input rather than as a score.

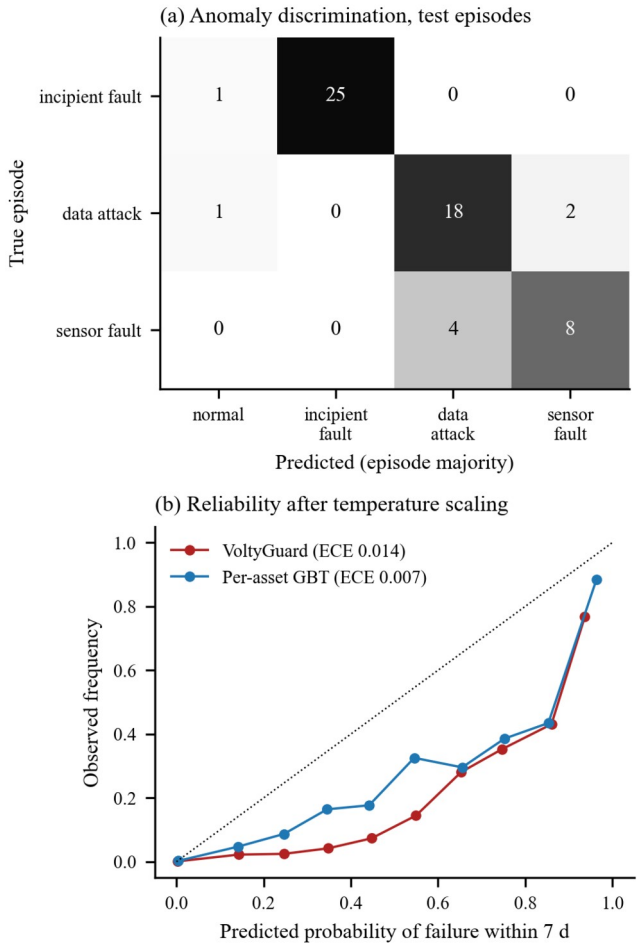


Fig. 4. (a) Episode-level confusion of the anomaly class head on the test realisations (seed 0). (b) Reliability diagram of the calibrated seven-day failure probability.

E. Ablation: What the Graph and the Standards Layers Add

On the three seeds shared by all networks (Fig. 5), removing the graph lowers recall from 70% to 63% at the F1 point and lowers recall from 83% to 80% at the matched budget, lowers precision from 88% to 79%, and makes the lead time more variable across seeds (s.d. 0.8 days against 1.1). Removing the standards and physics features while keeping the graph raises recall from 70% to 74% and lowers precision from 88% to 75%, lowers fault-versus-attack accuracy to 80%, and loses the mask attacks entirely (0 of 4 flagged, against 3 of 4 with the full model): without the hot-spot temperature implied by load, a replayed condition record is indistinguishable from a healthy one. The two layers are complementary. The physics features carry the thermal history that no 24-hour window of raw data contains and the reference against which a spoofed sensor is exposed; the graph carries the neighbour consistency that no single-asset model can test.

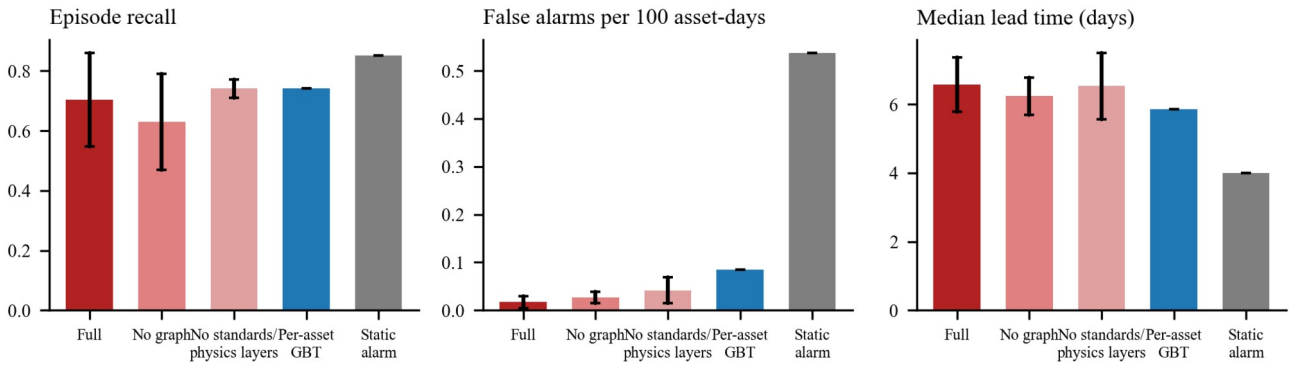


Fig. 5. Ablation on the test realisations (mean ± s.d. over three seeds): episode recall, false clusters per 100 asset-days and median warning lead time.

F. Regulation-Aware Actions and N-1 Feasibility

TABLE V
SERA ACTION MAPPING OF THE 27 TEST FAILURES (SEED 0); LAST ROW: ALERTS HELD AS SUSPECTED ATTACKS ARE NOT SCHEDULED

Model	Planned	Transfer	Emergency	Missed	N-1 feas.
Static alarm	18	2	3	4	2/5
Per-asset GBT	17	2	1	7	2/3
VoltyGuard	22	0	0	5	0/0
VoltyGuard, gated	22	0	0	5	0/0
VoltyGuard, gated + policy	21	—	0	5 (+1 held)	—

Table V applies the scheduler. With VoltyGuard gated, 22 of 27 failures can be handled as planned outages with the two-day customer notice, 0 as N-1 transfers without interruption and 0 require an emergency response; 5 are missed. The standards alarm alone leaves 7 failures to the emergency path. Of the 0 short-lead cases checked in the power-flow model, 0 satisfied the voltage and ampacity limits after tie-switch closure, a reminder that the value of a warning depends on the network's reserve as much as on its timing. Under the cyber-aware policy the gated model still schedules 21 failures as planned outages; 1 failure whose only alerts were classed as attacks would wait for the security review.

VI. Discussion

Three findings carry over to deployment. First, the standards layer is not a formality. On normal node-hours it raises caution in only 3.6% of cases, and gating the learned alerts on it costs no recall while giving every critical alert a reference an engineer can check against C57.104 or IEC 60270. Second, the gains from the graph are largest exactly where the attack surface is: the parent branch's power-balance residual is invisible to a per-asset model, and it is what separates a stealthy electrical injection from a genuine overload. Third, lead time is not a free variable. Transformer ageing is slow and observable, so warnings of one to several weeks are realistic; joint failures preceded by self-clearing faults give days, and only where a PD monitor or a fault recorder covers the segment. The scheduler makes this difference operational rather than statistical.

The study has limits that we state plainly. All data are synthetic, generated from standard models and documented assumptions; the degradation laws, gas generation rates and attack models are plausible, not fitted to Saudi field records, and the absolute numbers will move when they are. The feeder is balanced and radial, without distributed generation, and the hourly resolution hides sub-cycle phenomena that synchrophasor or fault-recorder data would expose. Model selection used four validation and six test realisations and 5 seeds; the seed-to-seed spread and the bootstrap intervals in Section V are the honest size of that uncertainty, and one seed's F1 threshold landed on the edge of the search grid and cost it the transformer failures, which argues for choosing operating points by false-alarm budget rather than by F1 in deployment. The next step is a shadow-mode pilot on a Saudi 13.8 kV network with read-only access to SCADA, condition-monitoring and relay data, which is the deployment path compatible with the National Cybersecurity Authority's operational-technology controls [28]: no control authority, data diode, engineer retains the decision, field outcomes fed back to retrain. Reliability effects would then be measured with the IEEE 1366 indices [29] rather than with simulated lead times, and cascading interactions [30] could be studied on a meshed model.

VII. Conclusion

Predictive maintenance becomes useful to a regulated operator when it is trusted, secure and convertible into admissible action. The framework evaluated here addresses all three conditions for Saudi MV networks: it gates learned predictions on diagnostic standards, it discriminates incipient faults from false data injection by exploiting physical and topological consistency, and it maps lead time to the actions SERA permits. On 24 simulated Riyadh summers of the IEEE 33-bus feeder it detected 82% of failures with a median warning of 7.2 days at the false-alarm budget of the standards alarm, separated faults from attacks in 90% of episodes, and, under a decision policy that holds suspected attacks for security review, turned most detected failures into planned outages while sending 0 of 1 alerted attacks to the scheduler. The generator, models and results are released for reuse, and the authors invite network operators to the shadow-mode pilot that must come next.

Data and Code Availability

The simulation generator, trained models, evaluation scripts and all result files are available at https://github.com/dr-elbehairy/voltyguard-mv-resilience (MIT licence).

References

[1] Daleel News, "Saudi Smart Grid Conference powers energy innovation," Dec. 19, 2025. [Online]. Available: https://news.daleel.gov.sa/energy/saudi-smart-grid-conference-powers-energy-innovation

[2] Attaqa, "Saudi Electricity Company 2024 results," Mar. 5, 2025. [Online]. Available: https://attaqa.net

[3] Saudi Electricity Regulatory Authority, "Electrical service level standards guide (guaranteed standards)," Umm Al-Qura Gazette, Jul. 2025. [Online]. Available: https://www.uqn.gov.sa/details?p=27249

[4] C. Rudin, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," Nat. Mach. Intell., vol. 1, no. 5, pp. 206–215, 2019, doi: 10.1038/s42256-019-0048-x.

[5] Y. Liu, P. Ning, and M. K. Reiter, "False data injection attacks against state estimation in electric power grids," in Proc. 16th ACM Conf. Comput. Commun. Secur. (CCS), Chicago, IL, USA, 2009, pp. 21–32, doi: 10.1145/1653662.1653666.

[6] IEEE Guide for the Interpretation of Gases Generated in Mineral Oil-Immersed Transformers, IEEE Std C57.104-2019, 2019, doi: 10.1109/IEEESTD.2019.8890040.

[7] Mineral oil-filled electrical equipment in service — Guidance on the interpretation of dissolved and free gases analysis, IEC 60599:2022.

[8] High-voltage test techniques — Partial discharge measurements, IEC 60270:2000+AMD1:2015.

[9] Power transformers — Part 7: Loading guide for mineral-oil-immersed power transformers, IEC 60076-7:2018.

[10] M. E. Baran and F. F. Wu, "Network reconfiguration in distribution systems for loss reduction and load balancing," IEEE Trans. Power Del., vol. 4, no. 2, pp. 1401–1407, 1989, doi: 10.1109/61.25627.

[11] M. Duval and A. dePablo, "Interpretation of gas-in-oil analysis using new IEC publication 60599 and IEC TC 10 databases," IEEE Electr. Insul. Mag., vol. 17, no. 2, pp. 31–41, 2001, doi: 10.1109/57.917529.

[12] D. Susa, M. Lehtonen, and H. Nordman, "Dynamic thermal modeling of distribution transformers," IEEE Trans. Power Del., vol. 20, no. 3, pp. 1919–1929, 2005, doi: 10.1109/TPWRD.2005.848675.

[13] B. D. Russell and C. L. Benner, "Intelligent systems for improved reliability and failure diagnosis in distribution systems," IEEE Trans.

Smart Grid, vol. 1, no. 1, pp. 48–56, 2010, doi: 10.1109/TSG.2010.2044898.

[14] T. S. Sidhu and Z. Xu, "Detection of incipient faults in distribution underground cables," IEEE Trans. Power Del., vol. 25, no. 3, pp. 1363–1371, 2010, doi: 10.1109/TPWRD.2010.2041373.

[15] S. Kulkarni, S. Santoso, and T. A. Short, "Incipient fault location algorithm for underground cables," IEEE Trans. Smart Grid, vol. 5, no. 3, pp. 1165–1174, 2014, doi: 10.1109/TSG.2014.2303483.

[16] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," in Proc. Int. Conf. Learn. Represent. (ICLR), Toulon, France, 2017.

[17] W. L. Hamilton, R. Ying, and J. Leskovec, "Inductive representation learning on large graphs," in Proc. Adv. Neural Inf. Process. Syst. (NeurIPS), vol. 30, Long Beach, CA, USA, 2017, pp. 1024–1034.

[18] W. Liao, B. Bak-Jensen, J. R. Pillai, Y. Wang, and Y. Wang, "A review of graph neural networks and their applications in power systems," J. Mod. Power Syst. Clean Energy, vol. 10, no. 2, pp. 345–360, 2022, doi: 10.35833/MPCE.2021.000058.

[19] K. Chen, J. Hu, Y. Zhang, Z. Yu, and J. He, "Fault location in power distribution systems via deep graph convolutional networks," IEEE J. Sel. Areas Commun., vol. 38, no. 1, pp. 119–131, 2020, doi: 10.1109/JSAC.2019.2951964.

[20] W. Li, D. Deka, M. Chertkov, and M. Wang, "Real-time faulted line localization and PMU placement in power systems through convolutional neural networks," IEEE Trans. Power Syst., vol. 34, no. 6, pp. 4640–4651, 2019, doi: 10.1109/TPWRS.2019.2917794.

[21] O. Kundacina, M. Cosovic, and D. Vukobratovic, "State estimation in electric power systems leveraging graph neural networks," in Proc. 17th Int. Conf. Probabilistic Methods Appl. Power Syst. (PMAPS), Manchester, U.K., 2022, pp. 1–6, doi: 10.1109/PMAPS53380.2022.9810559.

[22] O. Boyaci, A. Umunnakwe, A. Sahu, M. R. Narimani, M. Ismail, K. R. Davis, and E. Serpedin, "Graph neural networks based detection of stealth false data injection attacks in smart grids," IEEE Syst. J., vol. 16, no. 2, pp. 2946–2957, 2022, doi: 10.1109/JSYST.2021.3109082.

[23] M. Atrigna, A. Buonanno, R. Carli, G. Cavone, P. Scarabaggio, M. Valenti, G. Graditi, and M. Dotoli, "A machine learning approach to fault prediction of power distribution grids under heatwaves," IEEE Trans. Ind. Appl., vol. 59, no. 4, pp. 4835–4845, 2023, doi: 10.1109/TIA.2023.3262230.

[24] IEEE Standard for Synchrophasor Measurements for Power Systems, IEEE Std C37.118.1-2011, 2011, doi: 10.1109/IEEESTD.2011.6111219.

[25] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in Proc. 34th Int. Conf. Mach. Learn. (ICML), Sydney, Australia, 2017, pp. 1321–1330.

[26] L. Thurner, A. Scheidler, F. Schäfer, J.-H. Menke, J. Dollichon, F. Meier, S. Meinecke, and M. Braun, "pandapower — An open-source Python tool for convenient modeling, analysis, and optimization of electric power systems," IEEE Trans. Power Syst., vol. 33, no. 6, pp. 6510–6521, 2018, doi: 10.1109/TPWRS.2018.2829021.

[27] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in Proc. 22nd ACM SIGKDD Int. Conf. Knowl. Discov. Data Min., San Francisco, CA, USA, 2016, pp. 785–794, doi: 10.1145/2939672.2939785.

[28] National Cybersecurity Authority, Operational Technology Cybersecurity Controls (OTCC-1:2022), Riyadh, Saudi Arabia, 2022.

[29] IEEE Guide for Electric Power Distribution Reliability Indices, IEEE Std 1366-2022, 2022, doi: 10.1109/IEEESTD.2022.9955492.

[30] I. Dobson, B. A. Carreras, V. E. Lynch, and D. E. Newman, "Complex systems analysis of series of blackouts: Cascading failure, critical points, and self-organization," Chaos, vol. 17, no. 2, p. 026103, 2007, doi: 10.1063/1.2737822.

BIOGRAPHIES



Moustafa Elbehairy is founder and Co-Chief Executive Officer of Nukhbaat Hulul Alamal (NHA), Riyadh, Saudi Arabia, and a co-founder of SystechHub, Germany. He has more than eighteen years of experience in artificial intelligence, data science and risk analytics for government and critical-infrastructure clients in Saudi Arabia and Egypt, and lectures in statistics and insurance at Suez Canal University, Egypt. His research interests include graph learning for infrastructure, AI governance and risk modelling.



Ahmad Abuseeni is founder and Chief Executive Officer of Nukhbaat Hulul Alamal (NHA), Riyadh, Saudi Arabia, and an associate professor of computer information systems. He has more than eighteen years of experience in data management and governance consulting, information technology and artificial intelligence governance, in addition to his service as an associate professor in the Department of Management Information Systems at colleges of business administration in Saudi government universities. He is an expert in assessing data-management maturity and the requirements of the Saudi National Data Index (NDI) and National AI Index (NAII).



Ahmed Abdulrahman Al-Jafin is an electrical engineer and the founder and General Manager of Energy Towers Company for Engineering, Saudi Arabia. He has held management positions at the Ministry of Energy and the Saudi Electricity Company. His interests include medium-voltage distribution networks and their protection.

Nasser Al-Maawi is an electrical engineer with United Specialists Co., Saudi Arabia. He has 10 years of experience in distribution-network projects and field instrumentation.

Amir Raouf is a consultant in operations management and research with Karaz Platform, Saudi Arabia. He has 15 years of experience in operations management and med AI. His interests include maintenance scheduling and optimisation.